

Dynamic-Precision Training of Deep Neural Network Models

Paul Estano

With Silviu I. Filip & Elisa Riccietti

November 15, 2022

Universite de Rennes 1 - IRISA - Inria

1. Context & Motivation
2. Neural Network Learning in Low/Mixed-Precision
3. Optimization in Dynamic-Precision


Context & Motivation

DNNs are (almost) everywhere

- *Deep Learning* (DL) has enabled many breakthroughs in recent years.




- The development and the use of these algorithms have a consequential impact on the environment:



(Taken from [Strubell et al., 2019])

DNNs are getting bigger

Training DNNs is computationally and memory expensive¹.




How can we make it more efficient and scalable?

¹<https://medium.com/riselab/ai-and-memory-wall-2cb4265cb0b8>

DNN Quantization

- **Idea:** Use less bits to represent numbers during DNN training.




Taken from [Horowitz et al., 2014]

- **How do we use smaller number formats without hurting performance?**

Neural Network Learning in Low/Mixed-Precision

Neural Network Learning




- We study a neural network composed of L layers h_ℓ :

$$h_\ell(X; W) = \phi_\ell(W_\ell h_{\ell-1}(X; W))$$

$W = \{W_\ell\}_{\ell=1}^L$ a set of weight matrices, $\{\phi_\ell\}_{\ell=1}^L$ a set of non-linear functions.

Neural Network Learning




- We study a neural network composed of L layers h_ℓ :

$$h_\ell(X; W) = \phi_\ell(W_\ell h_{\ell-1}(X; W))$$

$W = \{W_\ell\}_{\ell=1}^L$ a set of weight matrices, $\{\phi_\ell\}_{\ell=1}^L$ a set of non-linear functions.

Neural Network Learning



- We study a neural network composed of L layers h_ℓ :


$$h_\ell(X; W) = \phi_\ell(W_\ell h_{\ell-1}(X; W))$$

$W = \{W_\ell\}_{\ell=1}^L$ a set of weight matrices, $\{\phi_\ell\}_{\ell=1}^L$ a set of non-linear functions.

- Given a training set $\{(X_i, Y_i)\}_{i=1}^N$ and an *error* function, we want to **minimize** the *empirical risk* f :

$$\arg \min_W f(W) = \arg \min_W \frac{1}{N} \sum_{i=1}^N \text{error}(h_L(X_i; W), Y_i)$$


The cost of Neural Network Training



$$W_{\ell}^{(k+1)} \leftarrow W_{\ell}^{(k)} - \gamma \frac{\partial f(W^{(k)})}{\partial W_{\ell}^{(k)}}$$

Requires several quantities:

The cost of Neural Network Training




$$W_{\ell}^{(k+1)} \leftarrow W_{\ell}^{(k)} - \gamma \frac{\partial f(W^{(k)})}{\partial W_{\ell}^{(k)}}$$

Requires several quantities:

- Forward activations (like for inference).

The cost of Neural Network Training




$$W_{\ell}^{(k+1)} \leftarrow W_{\ell}^{(k)} - \gamma \frac{\partial f(W^{(k)})}{\partial W_{\ell}^{(k)}}$$

Requires several quantities:

- Forward activations (like for inference).
- Backward gradients (*backpropagation*):
 - With respect to activations $\Delta_{h_{\ell}}$
 - With respect to weights $\Delta_{W_{\ell}}$

The cost of Neural Network Training




$$W_{\ell}^{(k+1)} \leftarrow W_{\ell}^{(k)} - \gamma \frac{\partial f(W^{(k)})}{\partial W_{\ell}^{(k)}}$$

Requires several quantities:

- ▶ Forward activations (like for inference).
- ▶ Backward gradients (*backpropagation*):
 - ▶ With respect to activations $\Delta_{h_{\ell}}$
 - ▶ With respect to weights $\Delta_{W_{\ell}}$

Can we use mixed-precision to reduce the cost of training?


Neural Network Learning in Low-Precision



- Operations are accessible up to a precision ϵ . We now look for:

$$\arg \min_W f(W) \text{ s.t. } |\hat{f}(W, \epsilon) - f(W)| \leq \epsilon$$


Neural Network Learning in Mixed-Precision



- Operations are accessible up to a precision ϵ_k . We now look for:

$$\arg \min_W f(W) \text{ s.t. } |\hat{f}(W, \epsilon_k) - f(W)| \leq \epsilon_k$$

Neural Network Learning in Mixed-Precision




- Operations are accessible up to a precision ϵ_k . We now look for:

$$\arg \min_W f(W) \text{ s.t. } |\hat{f}(W, \epsilon_k) - f(W)| \leq \epsilon_k$$

- Reducing precision during training is hard:
 - vanishing & exploding gradients
 - small parameter updates

Neural Network Learning in Mixed-Precision



- Operations are accessible up to a precision ϵ_k . We now look for:

$$\arg \min_W f(W) \text{ s.t. } |\hat{f}(W, \epsilon_k) - f(W)| \leq \epsilon_k$$

- Reducing precision during training is hard:
 - vanishing & exploding gradients
 - small parameter updates
- Large dynamic range required → use floating-point

Recent results in MP Training

Comparison of recent MP training methods:

Method	Formats ((Exponent, Mantissa) / Width)				Top-1 Accuracy	
	w	Δ_W	Δ_{he}	Acc.	FP32	Proposed
SWALP [Yang et al., 2019]	8	8	8	32	70.3	65.8
S2FP8 [Cambier et al., 2020]	(5,2)/(8,23)	(5,2)	(5,2)	(8,23)	70.3	69.6
HFP8 [Sun et al., 2019]	(4,3)	(6,9)	(5,2)	(6,9)	69.4	69.4
BM8 [Fox et al., 2021]	(2,5)	(6,9)	(4,3)	31	69.7	69.8
FP8-SEB [Park et al., 2022]	(4,3)	(4,3)	(4,3)	(8,23)	69.7	69.0
FP134 [Lee et al., 2022]	(3,4)	(3,4)	(3,4)	(8,23)	69.8	69.8

Results are ImageNet accuracy (%) using ResNet18 (adapted from [Lee et al., 2022]).

Notable Ideas:

- Shared scaling factor/bias at block or tensor level
→ shift dynamic range at runtime

Recent results in MP Training

Comparison of recent MP training methods:

Method	Formats ((Exponent, Mantissa) / Width)				Top-1 Accuracy	
	w	Δ_W	Δ_{he}	Acc.	FP32	Proposed
SWALP [Yang et al., 2019]	8	8	8	32	70.3	65.8
S2FP8 [Cambier et al., 2020]	(5,2)/(8,23)	(5,2)	(5,2)	(8,23)	70.3	69.6
HFP8 [Sun et al., 2019]	(4,3)	(6,9)	(5,2)	(6,9)	69.4	69.4
BM8 [Fox et al., 2021]	(2,5)	(6,9)	(4,3)	31	69.7	69.8
FP8-SEB [Park et al., 2022]	(4,3)	(4,3)	(4,3)	(8,23)	69.7	69.0
FP134 [Lee et al., 2022]	(3,4)	(3,4)	(3,4)	(8,23)	69.8	69.8

Results are ImageNet accuracy (%) using ResNet18 (adapted from [Lee et al., 2022]).

Notable Ideas:

- ▶ Shared scaling factor/bias at block or tensor level
→ shift dynamic range at runtime
- ▶ Scale the loss function before back propagation
→ shifts gradients in a representable range to avoid under/overflows

Recent results in MP Training

Comparison of recent MP training methods:

Method	Formats ((Exponent, Mantissa) / Width)				Top-1 Accuracy	
	w	Δ_W	Δ_{he}	Acc.	FP32	Proposed
SWALP [Yang et al., 2019]	8	8	8	32	70.3	65.8
S2FP8 [Cambier et al., 2020]	(5,2)/(8,23)	(5,2)	(5,2)	(8,23)	70.3	69.6
HFP8 [Sun et al., 2019]	(4,3)	(6,9)	(5,2)	(6,9)	69.4	69.4
BM8 [Fox et al., 2021]	(2,5)	(6,9)	(4,3)	31	69.7	69.8
FP8-SEB [Park et al., 2022]	(4,3)	(4,3)	(4,3)	(8,23)	69.7	69.0
FP134 [Lee et al., 2022]	(3,4)	(3,4)	(3,4)	(8,23)	69.8	69.8

Results are ImageNet accuracy (%) using ResNet18 (adapted from [Lee et al., 2022]).

Notable Ideas:

- ▶ Shared scaling factor/bias at block or tensor level
→ shift dynamic range at runtime
- ▶ Scale the loss function before back propagation
→ shifts gradients in a representable range to avoid under/overflows
- ▶ Ad-hoc rounding used in the quantizer: e.g. stochastic [Yang et al., 2019] & hysteresis [Lee et al., 2022]
- ▶ Most of these papers do not use any ad-hoc optimization method!

- Use of *QPytorch* [Zhang et al., 2019] → not an adequate simulation of low precision hardware.
Goal: Use *MPtorch* [Tatsumi et al., 2021]. Work in progress based on Pytorch and CUDA to simulate low-precision/mixed-precision computation.

- ▶ Use of *QPytorch* [Zhang et al., 2019] → not an adequate simulation of low precision hardware.
Goal: Use MPtorch [Tatsumi et al., 2021]. Work in progress based on Pytorch and CUDA to simulate low-precision/mixed-precision computation.
- ▶ Lack of theoretical results and use of heuristics.
Goal: Have convergence guarantees in low/mixed-precision.

Optimization in Dynamic-Precision

Optimization in dynamic precision

- ▶ **Goal:** a method with
 - ▶ an **automatic** strategy to set the precision
 - ▶ good **convergence properties**

Optimization in dynamic precision

- ▶ **Goal:** a method with
 - ▶ an **automatic** strategy to set the precision
 - ▶ good **convergence properties**

Optimization in dynamic precision

- ▶ **Goal:** a method with
 - ▶ an **automatic** strategy to set the precision
 - ▶ good **convergence properties**
- ▶ [Conn et al., 2000, Ch. 10.6] introduces a *trust-region* based algorithm satisfying such requirements in a **classical setting**:

$$\min_x f(x)$$

Optimization in dynamic precision

- ▶ **Goal:** a method with
 - ▶ an **automatic** strategy to set the precision
 - ▶ good **convergence properties**
- ▶ [Conn et al., 2000, Ch. 10.6] introduces a *trust-region* based algorithm satisfying such requirements in a **classical setting**:

$$\min_x f(x)$$

- ▶ We investigate the extension of this scheme to a **stochastic setting**:

$$\min_x \mathbb{E} [f(x)]$$

Optimization in dynamic precision

- ▶ **Goal:** a method with
 - ▶ an **automatic** strategy to set the precision
 - ▶ good **convergence properties**
- ▶ [Conn et al., 2000, Ch. 10.6] introduces a *trust-region* based algorithm satisfying such requirements in a **classical setting**:


$$\min_x f(x)$$

- ▶ We investigate the extension of this scheme to a **stochastic setting**:

$$\min_x \mathbb{E}[f(x)]$$

- ▶ Notable application: training of deep neural networks


Trust-Region (TR) Methods [Nocedal et al., 2006]



At every iteration:

- Generate a step p_k with the help of a model m_k of the loss f
- The model is *trusted* inside of a region of radius Δ_k
- Δ_k varies based on the model performance in previous iterations

Trust-Region (TR) Methods [Nocedal et al., 2006]




With $f_k = f(x_k)$, g_k and B_k resp. gradient and hessian of f

- TR methods often minimize a *quadratic model*

$$p_k = \arg \min_{\|p\| \leq \Delta_k} m_k(p) = \arg \min_{\|p\| \leq \Delta_k} f_k + g_k^T p + \frac{1}{2} p^T B_k p$$

Trust-Region (TR) Methods [Nocedal et al., 2006]




With $f_k = f(x_k)$, g_k and B_k resp. gradient and hessian of f

- TR methods often minimize a *quadratic model*, but using a *linear model* also works

$$p_k = \arg \min_{\|p\| \leq \Delta_k} m_k(p) = \arg \min_{\|p\| \leq \Delta_k} f_k + g_k^T p$$

Trust-Region (TR) Methods [Nocedal et al., 2006]



With $f_k = f(x_k)$, g_k and B_k resp. gradient and hessian of f

- TR methods often minimize a *quadratic model*, but using a *linear model* also works

$$p_k = \arg \min_{\|p\| \leq \Delta_k} m_k(p) = \arg \min_{\|p\| \leq \Delta_k} f_k + g_k^T p + \frac{1}{2} p^T B_k p$$

Basic Trust-Region

$$\min_x f(x)$$

Basic Trust-Region Algorithm [Conn et al., 2000, Sec. 6.1]

for $k = 0, 1, 2, \dots$ **do**

1. Compute model $m_k(p)$
2. Compute the step p_k solving the sub-problem

$$\arg \min_{\|p\| \leq \Delta_k} m_k(p)$$

$$3. \text{ Trust-region update: } \rho_k = \frac{\overbrace{f(x_k) - f(x_k + p_k)}^{\text{actual reduction}}}{\underbrace{m_k(0) - m_k(p_k)}_{\text{predicted reduction}}}$$

If $\rho_k \geq \eta_2$	Then $\Delta_{k+1} \leftarrow 2\Delta_k$	Accept $x_k + p_k$
If $\eta_2 > \rho_k \geq \eta_1$	Then $\Delta_{k+1} \leftarrow \Delta_k$	Accept $x_k + p_k$
If $\rho_k < \eta_1$	Then $\Delta_{k+1} \leftarrow 0.5\Delta_k$	Reject $x_k + p_k$

end for

where η_1, η_2 are constants satisfying $0 < \eta_1 \leq \eta_2 < 1$

Dynamic-Precision Trust-Region [Conn et al., 2000, Ch. 10.6]

- Operations are accessible up to a precision ϵ_k :

$$|\hat{f}(x, \epsilon_k) - f(x)| \leq \epsilon_k$$

- Want to solve $\min_x f(x)$ by having access just to $\hat{f}(x, \epsilon_k)$

Dynamic-Precision Trust-Region [Conn et al., 2000, Ch. 10.6]

- Operations are accessible up to a precision ϵ_k :

$$|\hat{f}(x, \epsilon_k) - f(x)| \leq \epsilon_k$$

- Want to solve $\min_x f(x)$ by having access just to $\hat{f}(x, \epsilon_k)$
- How can we choose the precision ϵ_k ?

Dynamic-Precision Trust-Region [Conn et al., 2000, Ch. 10.6]

- Operations are accessible up to a precision ϵ_k :

$$|\hat{f}(x, \epsilon_k) - f(x)| \leq \epsilon_k$$

- Want to solve $\min_x f(x)$ by having access just to $\hat{f}(x, \epsilon_k)$
- How can we choose the precision ϵ_k ?
 - guarantee monotonic decrease of the objective function f

Dynamic-Precision Trust-Region [Conn et al., 2000, Ch. 10.6]

- Operations are accessible up to a precision ϵ_k :

$$|\hat{f}(x, \epsilon_k) - f(x)| \leq \epsilon_k$$

- Want to solve $\min_x f(x)$ by having access just to $\hat{f}(x, \epsilon_k)$
- How can we choose the precision ϵ_k ?
 - guarantee monotonic decrease of the objective function f

Theory (intuition)

In mixed-precision $\rho_k \geq \eta_1$ gives us:

$$\hat{f}(x_k, \epsilon_k) - \hat{f}(x_k + p_k, \epsilon_k) \geq \eta_1 [m_k(0) - m_k(p_k)] > 0$$

Dynamic-Precision Trust-Region [Conn et al., 2000, Ch. 10.6]

- Operations are accessible up to a precision ϵ_k :

$$|\hat{f}(x, \epsilon_k) - f(x)| \leq \epsilon_k$$

- Want to solve $\min_x f(x)$ by having access just to $\hat{f}(x, \epsilon_k)$
- How can we choose the precision ϵ_k ?
 - guarantee monotonic decrease of the objective function f

Theory (intuition)

In mixed-precision $\rho_k \geq \eta_1$ gives us:

$$\hat{f}(x_k, \epsilon_k) - \hat{f}(x_k + p_k, \epsilon_k) \geq \eta_1 [m_k(0) - m_k(p_k)] > 0$$

By definition of \hat{f} :

$$\begin{aligned} f(x_k) - f(x_k + p_k) &\geq \hat{f}(x_k, \epsilon_k) - \hat{f}(x_k + p_k, \epsilon_k) - 2\epsilon_k \\ &\geq \underbrace{\eta_1 [m_k(0) - m_k(p_k)]}_{>0} - \underbrace{2\epsilon_k}_{?} > 0 \end{aligned}$$

Dynamic-Precision Trust-Region [Conn et al., 2000, Ch. 10.6]

- Operations are accessible up to a precision ϵ_k :

$$|\hat{f}(x, \epsilon_k) - f(x)| \leq \epsilon_k$$

- Want to solve $\min_x f(x)$ by having access just to $\hat{f}(x, \epsilon_k)$
- How can we choose the precision ϵ_k ?
 - guarantee monotonic decrease of the objective function f

Theory (intuition)

In mixed-precision $\rho_k \geq \eta_1$ gives us:

$$\hat{f}(x_k, \epsilon_k) - \hat{f}(x_k + p_k, \epsilon_k) \geq \eta_1 [m_k(0) - m_k(p_k)] > 0$$

By definition of \hat{f} :

$$\begin{aligned} f(x_k) - f(x_k + p_k) &\geq \hat{f}(x_k, \epsilon_k) - \hat{f}(x_k + p_k, \epsilon_k) - 2\epsilon_k \\ &\geq \underbrace{\eta_1 [m_k(0) - m_k(p_k)]}_{>0} - \underbrace{2\epsilon_k}_{?} > 0 \end{aligned}$$

$$\implies \epsilon_k < \eta_1 \frac{1}{2} [m_k(0) - m_k(p_k)]$$

Dynamic-Precision Trust-Region [Conn et al., 2000, Ch. 10.6]

$$\min_x f(x) \text{ s.t. } |\hat{f}(x, \epsilon_k) - f(x)| \leq \epsilon_k$$

Trust-Region with Dynamic-Precision

for $k = 0, 1, 2, \dots$ **do**

1. Compute model $m_k(p)$
2. Compute the step p_k solving the sub-problem

$$\arg \min_{\|p\| \leq \Delta_k} m_k(p)$$

if $\eta_1 \frac{1}{2} [m_k(0) - m_k(p_k)] > \epsilon_k$ **then**

3. Trust-region update: $\rho_k = \frac{\hat{f}(x_k, \epsilon_k) - \hat{f}(x_k + p_k, \epsilon_k)}{m_k(0) - m_k(p_k)}$

If $\rho_k \geq \eta_2$ Then $\Delta_{k+1} \leftarrow 2\Delta_k$ Accept $x_k + p_k$ & $\epsilon_{k+1} = \epsilon_k$

If $\eta_2 > \rho_k \geq \eta_1$ Then $\Delta_{k+1} \leftarrow \Delta_k$ Accept $x_k + p_k$ & $\epsilon_{k+1} = \epsilon_k$

If $\rho_k < \eta_1$ Then $\Delta_{k+1} \leftarrow 0.5\Delta_k$ Reject $x_k + p_k$ & $\epsilon_{k+1} = \epsilon_k$

else

Reduce ϵ_k and go to 1.

end if

end for

Dynamic-Precision Trust-Region

Limitations of the classical setting:

- ▶ small dimension: use a quadratic model
- ▶ deterministic setting: imposes a monotonic decrease of the loss

We adapt the previous algorithm to a stochastic setting:

- ▶ use batch training: stochasticity
- ▶ use a linear model: no hessian
- ▶ consider a non-monotonous loss function: ρ evaluated F times/epoch

Can be fine tuned for the training of neural networks:

- ▶ scale the loss: avoid vanishing gradients
- ▶ introduce momentum and weight-decay


$$\min_x f(x) + \lambda \|x\|_2^2 \text{ s.t. } |\hat{f}(x, \epsilon_k) - f(x)| \leq \epsilon_k$$

$$\text{Gradient estimate is } G_{k+1} = G_k + \mu \frac{\partial f(x)}{\partial x} \quad \& \quad G_0 = \frac{\partial f(x)}{\partial x}$$


Setup:

- ▶ ResNet-20 [Zhang et al., 2015]
- ▶ On CIFAR-10 (data augmented)
- ▶ 200 Epochs with batch size $B = 128$
- ▶ **TR**: $TR_{init} = 0.1$
- ▶ **TR**: Floating Point on 8, 16, 32 bits (E4M3, E5M10, IEEE 754)
- ▶ **SGD**: $LR_{init} = 0.1$ with cosine annealing scheduler
- ▶ **SGD**: Floating Point on 8 bits
- ▶ **TR&SGD**: Loss Scaling is used ($f \times 1024$)
- ▶ **TR&SGD**: Quantization with *QPytorch* [Zhang et al., 2019]

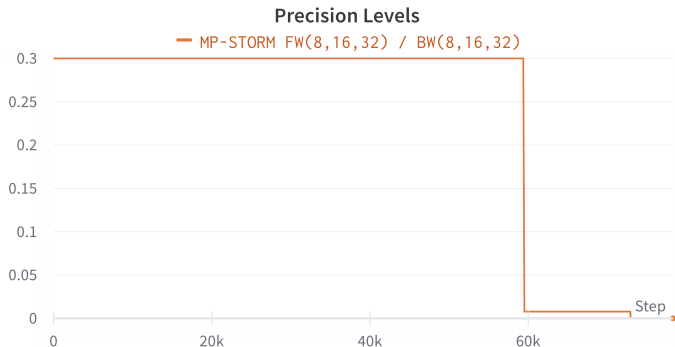
Initial Results: Test Errors



Initial Results: Loss



Initial Results: Precision



Limitations

- ▶ Convergence theory unclear (yet).
- ▶ Requires 2 evaluation of the loss/iteration.
- ▶ Mixed-precision does not seem to have an effect.
- ▶ TR methods are cumbersome

Limitations

- ▶ Convergence theory unclear (yet).
- ▶ Requires 2 evaluation of the loss/iteration.
- ▶ Mixed-precision does not seem to have an effect.
- ▶ TR methods are cumbersome
 - **Could we extend simpler (first-order) methods instead?**

Limitations

- ▶ Convergence theory unclear (yet).
- ▶ Requires 2 evaluation of the loss/iteration.
- ▶ Mixed-precision does not seem to have an effect.
- ▶ TR methods are cumbersome
 - **Could we extend simpler (first-order) methods instead?**
 - **Idea:** extend SGD to a mixed-precision setting

Mixed-Precision training of DNNs: Summary & Current Work

- ▶ SOTA methods usually use **heuristics** or focus on **quantization** aspects.
- ▶ We extended a classical optimization algorithms:
 - ▶ A **trust-region** method allowing us to jointly tune the step-size and the precision.

Current work:

- ▶ Running experiments:
 - ▶ Trust-Region with momentum
 - ▶ Trust-Region with weight-decay
 - ▶ Trust-Region with momentum and weight-decay
- ▶ Convergence of Stochastic Trust-Region
- ▶ Calculation error of a multi-layer perceptron
- ▶ Stochastic Gradient Descent in mixed-precision

Thank you! Questions?

Backup

Mixed-Precision Stochastic Gradient Descent: reminders 1

We extend SGD convergence proof [Ghadimi et al., 2013] to a mixed-precision problem

Definition

- We want to solve:

$$f^* = \inf_{x \in \mathbb{R}^n} f(x)$$

With f potentially non-convex

- We only have access to \hat{f} such that:

$$|\hat{f}(x_k, \epsilon_k) - f(x_k)| \leq \epsilon_k$$

- We only have access to a *stochastic gradient* $G(x_k, \xi_k)$ where ξ_k is a random variable in \mathbb{R}^d

Mixed-Precision Stochastic Gradient Descent: reminders 2

Assumption

- f is L -smooth:

$$\|\nabla f(x) - \nabla f(y)\| \leq L\|x - y\|.$$

- $G(x_k, \xi_k)$ is an unbiased estimator of $\nabla f(x_k)$ with bounded variance:

$$\mathbb{E}[G(x_k, \xi_k)] = \nabla f(x_k)$$

$$\mathbb{E}[\|G(x_k, \xi_k) - \nabla f(x_k)\|^2] \leq \sigma^2$$

- The precision ϵ_k , the step-size γ_k and the probability mass function $P_R(k)$ are such that $\gamma_k \leq \frac{2(1-\eta)}{L}$ and:

$$\text{Prob}\{R = k\} = \frac{2(1-\eta)\gamma_k - L\gamma_k^2}{\sum_{k=1}^N 2(1-\eta)\gamma_k - L\gamma_k^2} \ \& \ \epsilon_k \leq \gamma_k \frac{\eta}{2} \|G(x_k, \xi_k)\|^2.$$

Theorem (Informal)

Under previous assumptions we have:

$$\mathbb{E}[\|\nabla f(x_R)\|^2] \leq \frac{\hat{f}(x_1, \epsilon_1) - \hat{f}^* - \epsilon_1 + L\sigma^2 \sum_{k=1}^N \gamma_k^2}{\sum_{k=1}^N 2(1 - \eta)\gamma_k - L\gamma_k^2} \leq O\left(\frac{1}{\sqrt{N}}\right),$$

where the expectation is taken with respect to R and $\xi_{[N]} = (\xi_1, \dots, \xi_N)$

MP-SGD convergence proof: a brief overview

Proof.

We extend [Ghadimi et al., 2013].

By definition of $\hat{f}(x_k)$, we have:

$$\hat{f}(x_{k+1}, \epsilon_{k+1}) - \hat{f}(x_k, \epsilon_k) \leq f(x_{k+1}) - f(x_k) + \epsilon_{k+1} + \epsilon_k$$

Using the assumptions on f & summing up:

$$\begin{aligned} \sum_{k=1}^N \left(\gamma_k - \frac{L}{2} \gamma_k^2 \right) \|\nabla f(x_k)\|^2 &\leq \hat{f}(x_1, \epsilon_1) - \hat{f}^* - \sum_{k=1}^N \left(\gamma_k - L \gamma_k^2 \right) \langle \nabla f(x_k), \delta_k \rangle \\ &\quad + \sum_{k=1}^N \frac{L}{2} \gamma_k^2 \|\delta_k\|^2 + \sum_{k=1}^N 2\epsilon_k - \epsilon_1 \end{aligned}$$

With $\delta_k = G(x_k, \xi_k) - \nabla f(x_k)$. The problematic term is $\sum_{k=1}^N 2\epsilon_k$.

Using the constraints on ϵ_k and $P_R(k)$ we find the final result

□

- ▶ The constraint is very restrictive and does not allow many iterations in low-precision → **Is there a way to relax this?** (Potentially decreasing SGD convergence rate?)
- ▶ We now need to tune the learning-rate → Could we extend some adaptive method (Adagrad, ADAM...)?

Neural Network Error Analysis

What are we looking for?

- We consider the computation of

$$h_\ell(x) = \phi_\ell(W_\ell h_{\ell-1}(x)) \in \mathbb{R}^{n_\ell}$$

with $h_0(x) = x \in \mathbb{R}^{n_0}$, where $\phi_\ell : \mathbb{R}^{n_\ell} \mapsto \mathbb{R}^{n_\ell}$ is an activation function and $W_\ell \in \mathbb{R}^{n_\ell \times n_{\ell-1}}$ is a matrix of weights.

- **Goal:** We are trying to model the error propagation through the network layers inside the forward path. Where ϵ_ℓ is the error committed at layer ℓ

Error analysis of DNNs: a theoretical result

Theorem (Informal)

$$\varepsilon_{\ell}^h = \kappa_{\phi_{\ell}}(v_{\ell}) \circ \kappa_{v_{\ell}} \circ (\varepsilon_{\ell}^W + \bar{\varepsilon}_{\ell-1}^h) + \varepsilon_{\ell}^{\phi},$$


With

- ▶ $\varepsilon_{\ell}^h, \varepsilon_{\ell}^W, \varepsilon_{\ell}^{\phi}, \bar{\varepsilon}_{\ell-1}^h$, various errors committed at layer ℓ and $\ell - 1$
- ▶ $\kappa_{\phi_{\ell}}$, condition number of activation function
- ▶ $\kappa_{v_{\ell}}$, condition number of the weights/inputs vector product
- ▶ The errors are amplified by the condition numbers of each layers.

Error analysis of DNNs: the intuition behind

★ PE Need some help with the def of ϕ condition number.

◇ ER Use tanh rather than sigmoid




Sigmoid activation function $\phi(x)$ and its condition number $\kappa_{\phi(x)}$

Idea: When $\|x\|$ big, the output error is small. When $\|x\|$ is small the output error is big.

Error analysis of DNNs: the intuition behind

★ PE Need some help with the def of ϕ condition number.

◇ ER Use tanh rather than sigmoid



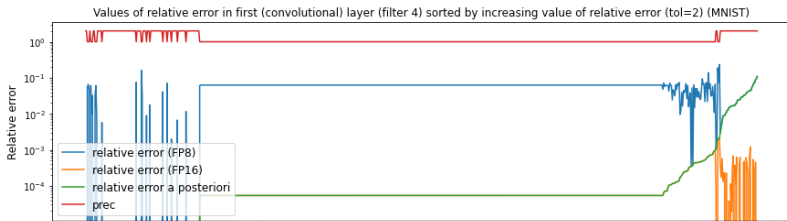
Sigmoid activation function $\phi(x)$ and its condition number $\kappa_{\phi(x)}$

Idea: When $\|x\|$ big, the output error is small. When $\|x\|$ is small the output error is big.

Our approach exploits this activation functions property to apply mixed-precision.


Error analysis of DNNs: Experimental validation 1

We select the precision at every dot product of every matrix multiplication based on a thresholding of the condition-numbers. **write the rule**




Error analysis of DNNs: Experimental validation 2

We vary the threshold (i.e. the ratio of FP8 operations) to observe its impact on the overall accuracy of a simple neural network (LeNet-5) for a common task (MNIST)



Error analysis of DNNs: Experimental validation 2

We vary the threshold (i.e. the ratio of FP8 operations) to observe its impact on the overall accuracy of a simple neural network (LeNet-5) for a common task (MNIST)



Future work: could we use this to accelerate DNNs inference?

Theory (intuition)

$$\begin{aligned}\hat{h}_\ell(x) &= \phi_\ell(v_\ell) \circ (\mathbf{1} + \kappa_{\phi_\ell}(v_\ell) \circ \delta v_\ell) \circ (\mathbf{1} + \Delta\phi_\ell) \\ &= h_\ell(x) \circ (\mathbf{1} + \underbrace{\kappa_{\phi_\ell}(v_\ell) \circ \delta v_\ell + \Delta\phi_\ell + \kappa_{\phi_\ell}(v_\ell) \circ \delta v_\ell \circ \Delta\phi_\ell}_{\Delta h_\ell})\end{aligned}$$

With δv_ℓ

$$m(W_k, s) = f_k + g(W_k)^T s$$

To refine precision ϵ_k of \hat{f} :

select


$$\epsilon_k^+ \in (0, \eta[m(W_k, 0) - m(W_k, s_k)] \text{ with } \eta > 0$$

i.e. for a linear model with $s_k = -\Delta_k \frac{g(W_k)}{\|g(W_k)\|_2}$

$$\epsilon_k^+ \in (0, -\eta \Delta_k \|g(W_k)\|_2]$$

Floating-Point 101

- Floating-point formats offer various trade-offs in terms of range, precision & performance




$$x = (-1)^s \times 1.m \times 2^{e - \text{BIAS}}$$

Format	Mantissa size	Exponent size	Bias	Range	Unit Roundoff	TFLOPS on H100	IEEE-754
fp64	52	11	1023	$10^{\pm 308}$	1×10^{-16}	48	Yes
fp32	23	8	127	$10^{\pm 38}$	1×10^{-8}	48	Yes
fp16	10	5	15	$10^{\pm 5}$	5×10^{-4}	400	Yes
tf32 (tf32)	10	8	127	$10^{\pm 38}$	5×10^{-4}	800	No
bfloat16 (bf16)	7	8	127	$10^{\pm 38}$	4×10^{-3}	800	No
fp8	3	4	7	$10^{\pm 2}$	6×10^{-2}	1600	No
	2	5	15	$10^{\pm 5}$	1×10^{-1}		

Floating-Point 101

- Floating-point formats offer various trade-offs in terms of range, precision & performance



$$x = (-1)^s \times 1.m \times 2^{e - \text{BIAS}}$$

Format	Mantissa size	Exponent size	Bias	Range	Unit Roundoff	TFLOPS on H100	IEEE-754
fp64	52	11	1023	$10^{\pm 308}$	1×10^{-16}	48	Yes
fp32	23	8	127	$10^{\pm 38}$	1×10^{-8}	48	Yes
fp16	10	5	15	$10^{\pm 5}$	5×10^{-4}	400	Yes
tf32 (tf32)	10	8	127	$10^{\pm 38}$	5×10^{-4}	800	No
bfloat16 (bf16)	7	8	127	$10^{\pm 38}$	4×10^{-3}	800	No
fp8	3	4	7	$10^{\pm 2}$	6×10^{-2}	1600	No
	2	5	15	$10^{\pm 5}$	1×10^{-1}		

- Default format for DNN applications is fp32