

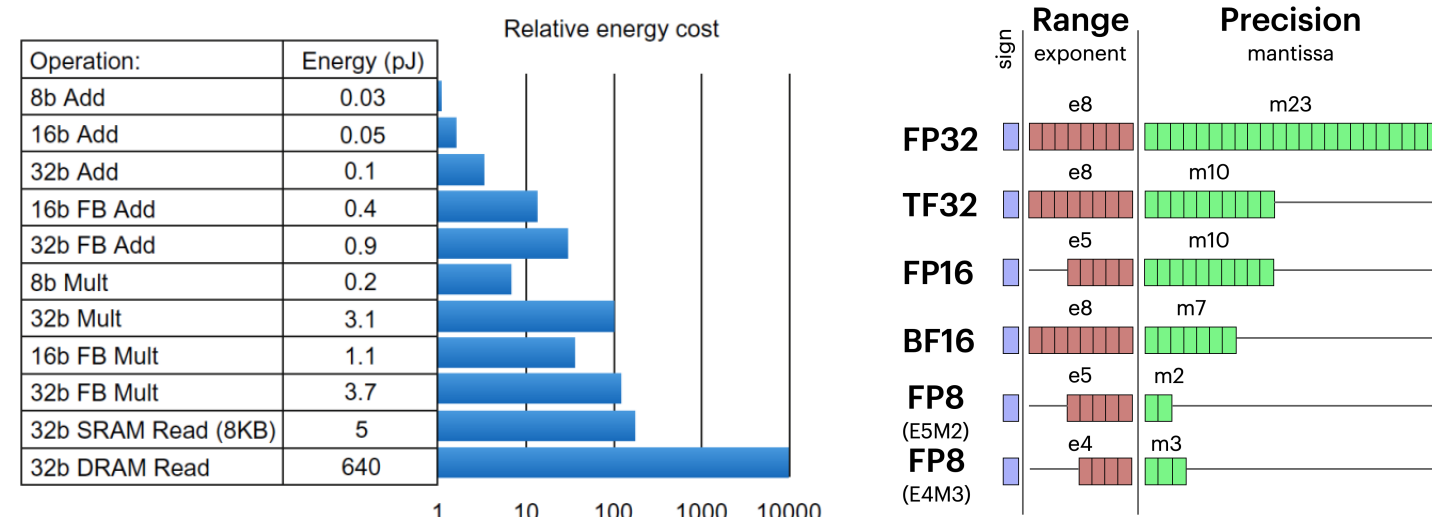
Dynamic Precision Training of Deep Neural Networks

Paul Estano Silviu I. Filip Elisa Riccietti

Universite de Rennes 1, IRISA, INRIA

Mixed Precision Learning

- **Objective:** make DNN training more efficient & scalable (for resource-constrained targets)
- Numerical values are representable using various formats (e.g., 32-bit, 16-bit, and 8-bit floating-point arithmetic)



[Horowitz et al., 2014]

- A small number format significantly reduces the **energy cost**
- Every function $f(x)$ is accessible up to a precision ϵ , through its low/mixed-precision approximation $\hat{f}(x, \epsilon)$:

$$|\hat{f}(x, \epsilon) - f(x)| \leq \epsilon$$

Mixed Precision Stochastic Gradient Descent (MP-SGD)

Extend Stochastic Gradient Descent by imposing a **constraint** on the precision.

$$\forall k \in \mathbb{N}^* : k \leq N, \epsilon_k \leq \eta \frac{\alpha_k}{2} \|\hat{g}(x_k, \epsilon_k)\|^2$$

$$x_{k+1} = x_k - \alpha_k \hat{g}(x_k, \epsilon_k)$$

With $\hat{g}(x_k, \epsilon_k)$ *stochastic gradient* up to a precision ϵ_k and $\eta < 1$

Theorem

With f possibly non-convex and L -smooth, MP-SGD converges to a stationary point. After N iterations we have:

$$\frac{1}{L} \mathbb{E}[\|\nabla f(x_R)\|^2] \leq O\left(\frac{1}{\sqrt{N}}\right)$$

where the expectation is taken with respect to R

Pros

- Theoretical **convergence guarantee**
- **Seamless** (no need to modify the update)
- Same complexity as SGD

Cons

- Requires strong assumptions on the **step-size** α_k
- Requires knowledge of L
- There are potentially better alternatives to SGD in practice
- Slow convergence rate

Intra-Layer Mixed Precision for Neural Networks Inference Acceleration

We derive a **bound** on the error propagated in the layers of a multi-layer perceptron due to mixed precision computation.

This bound depends on the **condition numbers** of the **weight matrices** and the **activation functions**.

We validate this bound through experimental results. In practice, we choose the precision based on the values of the condition numbers.

Layer-wise Error Analysis

The error committed at layer ℓ is bounded by ϵ_ℓ such that:

$$\epsilon_\ell^h = \kappa_{\phi_\ell}(v_\ell) \circ \kappa_{v_\ell} \circ (\epsilon_\ell^W + \bar{\epsilon}_{\ell-1}^h) + \epsilon_\ell^\phi,$$

where:


- κ_{ϕ_ℓ} condition number of the activation function
- κ_{v_ℓ} condition number of the weights/inputs vector product
- $\epsilon_\ell^h, \epsilon_\ell^W, \epsilon_\ell^\phi, \bar{\epsilon}_{\ell-1}^h$ various errors committed at layers ℓ and $\ell - 1$

A Posteriori Mixed Precision Linear Layer Computation


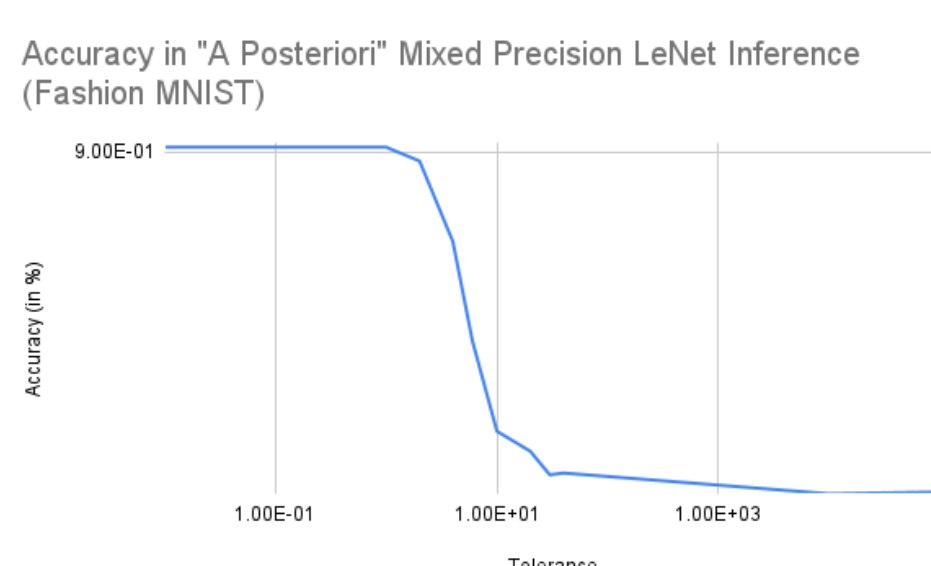
1. Compute the layer output in full precision
 2. Compute condition numbers κ_{ϕ_ℓ} & κ_{v_ℓ} and their product in full precision
- for every element κ_i in the vector $\kappa_{\phi_\ell} \circ \kappa_{v_\ell}$ do**
- if $\kappa_i \leq \text{tolerance}$ then**
- $\epsilon_{\ell,i}^h \leftarrow \text{FP8}$
- else**
- $\epsilon_{\ell,i}^h \leftarrow \text{FP16}$
- end if**
- end for**
3. Compute the layer using ϵ_ℓ^h values as precision for every dot product in the layer

Experiments on MNIST and Fashion-MNIST with LeNet-5

Selecting the precision using the a posteriori algorithm significantly reduces the relative error at each layer and globally.




The error (to a large extent) depends on the **activation functions condition numbers** \Rightarrow precision of a given network can be chosen **independently from the data**.



Neural Networks Training in **Trust**-able Precision

- Trusts and optimizes a simple **model** m , $m_k(p) = f_k + g_k^T p$ within a ball
- The step-size is automatically chosen
- Extend [Gratton et al., 2019]
- Uses a **stochastic** setting and a **1st order** model
- Same complexity as SGD



MP-STORM Algorithm

- for** $k = 1, 2, \dots, N$ **do**
1. Compute the step p_k
 2. Set the precision depending on p_k
 3. Update p_k with the new precision if required
 4. Accept or reject the step depending on p_k
 5. Update the *trust-radius* Δ_k depending on p_k
- end for**

$$\rho_k = \frac{\overbrace{f(x_k) - f(x_k + p_k)}^{\text{actual reduction}}}{\underbrace{m_k(0) - m_k(p_k)}_{\text{predicted reduction}}}$$

CIFAR-10 Experiments

Encouraging results on CIFAR-10, competitive with SOTA FP32 methods.

